

```

import numpy as np      # import numpy library
import pandas as pd     # import pandas library for accessing and
                        # analyzing the data
from sklearn.impute import KNNImputer
#KNN Iputation library for handling missing data commented out after
#processing once and stored the imputed data in new file as it takes 1
#hour to process,
from sklearn.preprocessing import LabelEncoder
label_encoder = LabelEncoder()
import matplotlib.pyplot as plt # import matplotlib library for plots
and visualization
import seaborn as sns
from sklearn.model_selection import train_test_split # import train-
test split for splitting the data into train and test
from sklearn.preprocessing import MinMaxScaler #library used for
scaling and standardizing the data
%matplotlib inline
#It is used to plot the matplotlib charts just below the code cells

df=pd.read_csv('Loan_Cleaned_10.csv')

```

df

| Score \ | Unnamed: 0           | Loan Status   | Current Loan Amount | Term | Credit |
|---------|----------------------|---------------|---------------------|------|--------|
| 0       | 0                    | 1             | 11520               | 1    |        |
| 741.0   |                      |               |                     |      |        |
| 1       | 1                    | 1             | 3441                | 1    |        |
| 734.0   |                      |               |                     |      |        |
| 2       | 2                    | 1             | 21029               | 1    |        |
| 747.0   |                      |               |                     |      |        |
| 3       | 3                    | 1             | 18743               | 1    |        |
| 747.0   |                      |               |                     |      |        |
| 4       | 4                    | 1             | 11731               | 1    |        |
| 746.0   |                      |               |                     |      |        |
| ...     | ...                  | ...           | ...                 | ...  | .      |
| ..      |                      |               |                     |      |        |
| 230998  | 256444               | 0             | 11953               | 1    |        |
| 717.0   |                      |               |                     |      |        |
| 230999  | 256446               | 1             | 3911                | 1    |        |
| 718.0   |                      |               |                     |      |        |
| 231000  | 256447               | 1             | 5078                | 1    |        |
| 737.0   |                      |               |                     |      |        |
| 231001  | 256448               | 0             | 12116               | 1    |        |
| 746.0   |                      |               |                     |      |        |
| 231002  | 256450               | 1             | 27902               | 0    |        |
| 678.0   |                      |               |                     |      |        |
|         | Years in current job | Annual Income | Monthly Debt        | \    |        |
| 0       | 10.0                 | 33694.0       | 584.03              |      |        |

|        |      |          |         |
|--------|------|----------|---------|
| 1      | 4.0  | 42269.0  | 1106.04 |
| 2      | 10.0 | 90126.0  | 1321.85 |
| 3      | 10.0 | 38072.0  | 751.92  |
| 4      | 4.0  | 50025.0  | 355.18  |
| ...    | ...  | ...      | ...     |
| 230998 | 10.0 | 39844.0  | 982.82  |
| 230999 | 2.0  | 90041.0  | 1706.58 |
| 231000 | 10.0 | 77186.0  | 1376.47 |
| 231001 | 9.0  | 52504.0  | 297.96  |
| 231002 | 10.0 | 117480.0 | 2111.38 |

|        | Years of Credit History | Months since last delinquent | ... | \ |
|--------|-------------------------|------------------------------|-----|---|
| 0      | 12.3                    | 41.0                         | ... |   |
| 1      | 26.3                    | 24.0                         | ... |   |
| 2      | 28.8                    | 35.6                         | ... |   |
| 3      | 26.2                    | 40.0                         | ... |   |
| 4      | 11.5                    | 42.4                         | ... |   |
| ...    | ...                     | ...                          | ... |   |
| 230998 | 11.7                    | 52.2                         | ... |   |
| 230999 | 19.9                    | 47.8                         | ... |   |
| 231000 | 19.1                    | 47.0                         | ... |   |
| 231001 | 15.1                    | 82.0                         | ... |   |
| 231002 | 18.0                    | 11.0                         | ... |   |

|        | Number of Credit Problems | Current Credit Balance | \ |
|--------|---------------------------|------------------------|---|
| 0      | 0                         | 6760.0                 |   |
| 1      | 0                         | 6262.0                 |   |
| 2      | 0                         | 20967.0                |   |
| 3      | 0                         | 22529.0                |   |
| 4      | 0                         | 17391.0                |   |
| ...    | ...                       | ...                    |   |
| 230998 | 1                         | 4176.0                 |   |
| 230999 | 0                         | 39804.7                |   |
| 231000 | 0                         | 1717.0                 |   |
| 231001 | 0                         | 3315.0                 |   |
| 231002 | 0                         | 28317.0                |   |

|            | Maximum Open Credit | Bankruptcies | Tax Liens | Home Mortgage |
|------------|---------------------|--------------|-----------|---------------|
| Own Home \ |                     |              |           |               |
| 0          | 16056               | 0            | 0         | 1             |
| 0          |                     |              |           |               |
| 1          | 19149               | 0            | 0         | 1             |
| 0          |                     |              |           |               |
| 2          | 28335               | 0            | 0         | 1             |
| 0          |                     |              |           |               |
| 3          | 43915               | 0            | 0         | 0             |
| 1          |                     |              |           |               |
| 4          | 37081               | 0            | 0         | 0             |
| 0          |                     |              |           |               |
| ...        | ...                 | ...          | ...       | ...           |

```

...
230998          4783          1          0          1
0
230999          44080          0          0          0
0
231000          9758          0          0          0
1
231001          20090          0          0          1
0
231002          62371          0          0          1
0

```

|        | Purpose | Unnamed: 19 | Unnamed: 20 |
|--------|---------|-------------|-------------|
| 0      | 203605  | NaN         | NaN         |
| 1      | 14196   | NaN         | NaN         |
| 2      | 203605  | NaN         | NaN         |
| 3      | 203605  | NaN         | NaN         |
| 4      | 203605  | NaN         | NaN         |
| ...    | ...     | ...         | ...         |
| 230998 | 203605  | NaN         | NaN         |
| 230999 | 203605  | NaN         | NaN         |
| 231000 | 203605  | NaN         | NaN         |
| 231001 | 203605  | NaN         | NaN         |
| 231002 | 203605  | NaN         | NaN         |

```
[231003 rows x 21 columns]
```

```
X=df.drop(columns=['Loan Status','Unnamed: 0','Unnamed: 19','Unnamed: 20'])
```

```
type(X)
```

```
pandas.core.frame.DataFrame
```

```
y = df['Loan Status']
```

```
y.value_counts()
```

```
Loan Status
```

```
1    138315
```

```
0     92688
```

```
Name: count, dtype: int64
```

```
# Scaling
```

```
scaler = MinMaxScaler()
```

```
X_scaled=scaler.fit_transform(X)
```

```
X = pd.DataFrame(X_scaled,columns=X.columns)
```

```
X_train, X_test, Y_train, Y_test = train_test_split(X, y, test_size = 0.25, random_state=100)
```

```
X_train.head(2)
```

|        | Current Loan Amount | Term | Credit Score | Years in current job |
|--------|---------------------|------|--------------|----------------------|
| \      |                     |      |              |                      |
| 199453 | 0.000239            | 0.0  | 0.653333     | 0.3                  |
| 97271  | 0.000042            | 1.0  | 0.346667     | 0.8                  |

|        | Annual Income | Monthly Debt | Years of Credit History | \ |
|--------|---------------|--------------|-------------------------|---|
| 199453 | 1.000000      | 0.692041     | 0.266766                |   |
| 97271  | 0.260097      | 0.162323     | 0.143070                |   |

|        | Months since last delinquent | Number of Open Accounts | \ |
|--------|------------------------------|-------------------------|---|
| 199453 | 0.172727                     | 0.236842                |   |
| 97271  | 0.125000                     | 0.197368                |   |

|        | Number of Credit Problems | Current Credit Balance | \ |
|--------|---------------------------|------------------------|---|
| 199453 | 0.0                       | 0.234762               |   |
| 97271  | 0.0                       | 0.212523               |   |

|          | Maximum Open Credit | Bankruptcies | Tax Liens | Home Mortgage |
|----------|---------------------|--------------|-----------|---------------|
| Own Home |                     |              |           |               |
| \        |                     |              |           |               |
| 199453   | 0.000184            | 0.0          | 0.0       | 1.0           |
| 0.0      |                     |              |           |               |
| 97271    | 0.000076            | 0.0          | 0.0       | 0.0           |
| 1.0      |                     |              |           |               |

|        | Purpose  |
|--------|----------|
| 199453 | 1.000000 |
| 97271  | 0.068548 |

```
X_train.shape
```

```
(173252, 17)
```

```
X_test.head(2)
```

|        | Current Loan Amount | Term | Credit Score | Years in current job |
|--------|---------------------|------|--------------|----------------------|
| \      |                     |      |              |                      |
| 196836 | 0.000054            | 1.0  | 0.848000     | 1.0                  |
| 61487  | 0.000046            | 1.0  | 0.405333     | 1.0                  |

|        | Annual Income | Monthly Debt | Years of Credit History | \ |
|--------|---------------|--------------|-------------------------|---|
| 196836 | 0.376166      | 0.119752     | 0.339791                |   |
| 61487  | 0.419985      | 0.096373     | 0.189270                |   |

|        | Months since last delinquent | Number of Open Accounts | \ |
|--------|------------------------------|-------------------------|---|
| 196836 | 0.264773                     | 0.105263                |   |

|       |          |          |
|-------|----------|----------|
| 61487 | 0.213636 | 0.078947 |
|-------|----------|----------|

|        | Number of Credit Problems | Current Credit Balance \ |
|--------|---------------------------|--------------------------|
| 196836 | 0.0                       | 0.317542                 |
| 61487  | 0.0                       | 0.238346                 |

|            | Maximum Open Credit | Bankruptcies | Tax Liens | Home Mortgage |
|------------|---------------------|--------------|-----------|---------------|
| Own Home \ |                     |              |           |               |
| 196836     | 0.000131            | 0.0          | 0.0       | 0.0           |
| 0.0        |                     |              |           |               |
| 61487      | 0.000138            | 0.0          | 0.0       | 1.0           |
| 0.0        |                     |              |           |               |

|        | Purpose  |
|--------|----------|
| 196836 | 1.000000 |
| 61487  | 0.014778 |

```
X_train.shape
(173252, 17)
X_train.isna().sum()
Current Loan Amount      0
Term                     0
Credit Score              0
Years in current job      0
Annual Income             0
Monthly Debt              0
Years of Credit History   0
Months since last delinquent 0
Number of Open Accounts   0
Number of Credit Problems 0
Current Credit Balance    0
Maximum Open Credit       0
Bankruptcies              0
Tax Liens                 0
Home Mortgage             0
Own Home                  0
Purpose                   0
dtype: int64

X_test.shape
(57751, 17)

Y_train.head(2)
199453    1
97271     1
Name: Loan Status, dtype: int64
```

```
Y_train.value_counts()
```

```
Loan Status
```

```
1    103748
```

```
0     69504
```

```
Name: count, dtype: int64
```

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score,
classification_report, ConfusionMatrixDisplay, \
                                precision_score, recall_score, f1_score,
roc_auc_score, roc_curve

models={
    "Logistic Regression":LogisticRegression(),
    "Decision Tree":DecisionTreeClassifier(),
    "Random Forest":RandomForestClassifier(),
    "Gradient Boost":GradientBoostingClassifier()
}
for i in range(len(list(models))):
    model = list(models.values())[i]
    model.fit(X_train, Y_train) # Train model

    # Make predictions
    y_train_pred = model.predict(X_train)
    y_test_pred = model.predict(X_test)

    # Training set performance
    model_train_accuracy = accuracy_score(Y_train, y_train_pred) #
    Calculate Accuracy
    model_train_f1 = f1_score(Y_train, y_train_pred,
average='weighted') # Calculate F1-score
    model_train_precision = precision_score(Y_train, y_train_pred) #
    Calculate Precision
    model_train_recall = recall_score(Y_train, y_train_pred) #
    Calculate Recall
    model_train_rocauc_score = roc_auc_score(Y_train, y_train_pred)

    # Test set performance
    model_test_accuracy = accuracy_score(Y_test, y_test_pred) #
    Calculate Accuracy
    model_test_f1 = f1_score(Y_test, y_test_pred, average='weighted')
    # Calculate F1-score
    model_test_precision = precision_score(Y_test, y_test_pred) #
    Calculate Precision
    model_test_recall = recall_score(Y_test, y_test_pred) # Calculate
```

Recall

```
model_test_rocauc_score = roc_auc_score(Y_test, y_test_pred)
#Calculate Roc
```

```
print(list(models.keys())[i])

print('Model performance for Training set')
print("- Accuracy: {:.4f}".format(model_train_accuracy))
print("- F1 score: {:.4f}".format(model_train_f1))

print("- Precision: {:.4f}".format(model_train_precision))
print("- Recall: {:.4f}".format(model_train_recall))
print("- Roc Auc Score: {:.4f}".format(model_train_rocauc_score))

print('-----')

print('Model performance for Test set')
print("- Accuracy: {:.4f}".format(model_test_accuracy))
print("- F1 score: {:.4f}".format(model_test_f1))
print("- Precision: {:.4f}".format(model_test_precision))
print("- Recall: {:.4f}".format(model_test_recall))
print("- Roc Auc Score: {:.4f}".format(model_test_rocauc_score))

print('='*35)
print('\n')
```

Logisitic Regression

Model performance for Training set

- Accuracy: 0.6285
- F1 score: 0.5857
- Precision: 0.6381
- Recall: 0.8771
- Roc Auc Score: 0.5672

-----

Model performance for Test set

- Accuracy: 0.6290
- F1 score: 0.5866
- Precision: 0.6384
- Recall: 0.8768
- Roc Auc Score: 0.5682

=====

Decision Tree

Model performance for Training set

- Accuracy: 1.0000

```

- F1 score: 1.0000
- Precision: 1.0000
- Recall: 1.0000
- Roc Auc Score: 1.0000
-----
Model performance for Test set
- Accuracy: 0.5552
- F1 score: 0.5561
- Precision: 0.6307
- Recall: 0.6196
- Roc Auc Score: 0.5393
=====

Random Forest
Model performance for Training set
- Accuracy: 1.0000
- F1 score: 1.0000
- Precision: 1.0000
- Recall: 1.0000
- Roc Auc Score: 1.0000
-----
Model performance for Test set
- Accuracy: 0.6380
- F1 score: 0.6187
- Precision: 0.6601
- Recall: 0.8149
- Roc Auc Score: 0.5946
=====

Gradient Boost
Model performance for Training set
- Accuracy: 0.6346
- F1 score: 0.5981
- Precision: 0.6449
- Recall: 0.8674
- Roc Auc Score: 0.5772
-----
Model performance for Test set
- Accuracy: 0.6317
- F1 score: 0.5948
- Precision: 0.6428
- Recall: 0.8657
- Roc Auc Score: 0.5743
=====

```



- Logistic Regression demonstrates consistent performance across training and test sets, with accuracy around 77%.
- It achieves a high recall (~96%), meaning it correctly identifies almost all positive cases (approved loans).
- The precision (~78%) indicates it manages false positives relatively well. However, the low ROC AUC (~0.55) suggests the model struggles to differentiate between approved and denied loans effectively.
- Inference: Reliable for moderately balanced datasets and interpretable results but limited in handling complex relationships.

### *Decision Tree*

- Decision Trees achieve perfect performance on the training set (100%), indicating overfitting to the training data.
- Test accuracy drops to ~69%, with a marginal improvement in ROC AUC (~0.59), showing reduced generalization to unseen data.
- While recall (~78%) is reasonable, the precision (~80%) indicates that the model is prone to false positives.
- Inference: Overfits easily; not ideal unless tuned to prevent memorization.

### *Random Forest*

- Random Forest mitigates the overfitting problem seen in Decision Trees by averaging across multiple trees.
- It achieves an accuracy of ~78% on the test set, with balanced recall (~96%) and precision (~80%).
- The ROC AUC (~0.59) still shows limited ability to distinguish between approved and denied loans.
- Inference: A robust model for handling large datasets and complex relationships but requires optimization for better performance.

### *Gradient Boosting*

- Gradient Boosting delivers similar accuracy (~77%) and recall (~97%) to Logistic Regression but at the cost of slightly reduced interpretability.
- Its ROC AUC (~0.55) suggests challenges in separating loan approval classes effectively. While it avoids overfitting seen in Decision Trees, its performance gains over Logistic Regression are minimal. -Inference: Performs well with non-linear data but may not offer significant advantages over simpler models in this scenario.

### *Final-Algo to Use*

- Random Forest

```
rf_params = {
    'n_estimators': [100, 200, 500],
    'max_depth': [10, 20, 30, None],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
```

```

    'bootstrap': [True, False]
}

gb_params = {
    'n_estimators': [100, 200, 500],
    'learning_rate': [0.01, 0.05, 0.1, 0.2],
    'max_depth': [3, 5, 7, 10],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'subsample': [0.8, 0.9, 1.0], # Fraction of samples to use for
    fitting each tree
    'criterion': ['friedman_mse', 'mse'] # Splitting criterion for
    tree building
}

```

```

random_cv_model = [
    ("RF", RandomForestClassifier(), rf_params),
    ("GB", GradientBoostingClassifier(), gb_params)
]

```

```

from sklearn.model_selection import RandomizedSearchCV

```

```

model_param = {}
for name, model, params in random_cv_model:
    random = RandomizedSearchCV(estimator=model,
                                param_distributions=params,
                                n_iter=20,
                                cv=5,
                                verbose=2,
                                n_jobs=-1)

    random.fit(X_train, Y_train)
    model_param[name] = random.best_params_

for model_name in model_param:
    print(f"----- Best Params for {model_name}")
    print(model_param[model_name])

```

Fitting 5 folds for each of 20 candidates, totalling 100 fits

c:\Users\prakhar\anaconda\Lib\site-packages\sklearn\model\_selection\\_validation.py:547: FitFailedWarning:

6 fits failed out of a total of 100.

The score on these train-test partitions for these parameters will be set to nan.

If these failures are not expected, you can try to debug them by setting error\_score='raise'.

Below are more details about the failures:

```

1 fits failed with the following error:
Traceback (most recent call last):
  File "c:\Users\prakhar\anaconda\Lib\site-packages\sklearn\
model_selection\_validation.py", line 895, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "c:\Users\prakhar\anaconda\Lib\site-packages\sklearn\base.py",
line 1474, in wrapper
    return fit_method(estimator, *args, **kwargs)
    ~~~~~
  File "c:\Users\prakhar\anaconda\Lib\site-packages\sklearn\ensemble\
_forest.py", line 489, in fit
    trees = Parallel(
    ~~~~~
  File "c:\Users\prakhar\anaconda\Lib\site-packages\sklearn\utils\
parallel.py", line 67, in __call__
    return super().__call__(iterable_with_config)
    ~~~~~
  File "c:\Users\prakhar\anaconda\Lib\site-packages\joblib\
parallel.py", line 1918, in __call__
    return output if self.return_generator else list(output)
    ~~~~~
  File "c:\Users\prakhar\anaconda\Lib\site-packages\joblib\
parallel.py", line 1847, in _get_sequential_output
    res = func(*args, **kwargs)
    ~~~~~
  File "c:\Users\prakhar\anaconda\Lib\site-packages\sklearn\utils\
parallel.py", line 129, in __call__
    return self.function(*args, **kwargs)
    ~~~~~
  File "c:\Users\prakhar\anaconda\Lib\site-packages\sklearn\ensemble\
_forest.py", line 200, in _parallel_build_trees
    tree.fit(
  File "c:\Users\prakhar\anaconda\Lib\site-packages\sklearn\tree\
_classes.py", line 305, in fit
    classes_k, y_encoded[:, k] = np.unique(y[:, k],
return_inverse=True)
~~~~~
  File "c:\Users\prakhar\anaconda\Lib\site-packages\numpy\lib\
arraysetops.py", line 274, in unique
    ret = _unique1d(ar, return_index, return_inverse, return_counts,
    ~~~~~
  File "c:\Users\prakhar\anaconda\Lib\site-packages\numpy\lib\
arraysetops.py", line 359, in _unique1d
    inv_idx = np.empty(mask.shape, dtype=np.intp)
    ~~~~~
numpy.core.exceptions.ArrayMemoryError: Unable to allocate 1.06 MiB

```

for an array with shape (138602,) and data type int64

```
-----  
-----  
1 fits failed with the following error:  
Traceback (most recent call last):  
  File "c:\Users\prakhar\anaconda\Lib\site-packages\sklearn\  
model_selection\_validation.py", line 895, in _fit_and_score  
    estimator.fit(X_train, y_train, **fit_params)  
  File "c:\Users\prakhar\anaconda\Lib\site-packages\sklearn\base.py",  
line 1474, in wrapper  
    return fit_method(estimator, *args, **kwargs)  
    ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^  
  File "c:\Users\prakhar\anaconda\Lib\site-packages\sklearn\ensemble\  
_forest.py", line 489, in fit  
    trees = Parallel(  
    ^^^^^^^  
  File "c:\Users\prakhar\anaconda\Lib\site-packages\sklearn\utils\  
parallel.py", line 67, in __call__  
    return super().__call__(iterable_with_config)  
    ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^  
  File "c:\Users\prakhar\anaconda\Lib\site-packages\joblib\  
parallel.py", line 1918, in __call__  
    return output if self.return_generator else list(output)  
    ^^^^^^^^^^^  
  File "c:\Users\prakhar\anaconda\Lib\site-packages\joblib\  
parallel.py", line 1847, in _get_sequential_output  
    res = func(*args, **kwargs)  
    ^^^^^^^^^^^^^^^^^^^^^^^  
  File "c:\Users\prakhar\anaconda\Lib\site-packages\sklearn\utils\  
parallel.py", line 129, in __call__  
    return self.function(*args, **kwargs)  
    ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^  
  File "c:\Users\prakhar\anaconda\Lib\site-packages\sklearn\ensemble\  
_forest.py", line 200, in _parallel_build_trees  
    tree._fit(  
  File "c:\Users\prakhar\anaconda\Lib\site-packages\sklearn\tree\  
_classes.py", line 472, in _fit  
    builder.build(self.tree_, X, y, sample_weight,  
missing_values_in_feature_mask)  
  File "sklearn\tree\_tree.pyx", line 166, in  
sklearn.tree._tree.DepthFirstTreeBuilder.build  
  File "sklearn\tree\_tree.pyx", line 285, in  
sklearn.tree._tree.DepthFirstTreeBuilder.build  
  File "sklearn\tree\_tree.pyx", line 940, in  
sklearn.tree._tree.Tree._add_node  
  File "sklearn\tree\_tree.pyx", line 908, in  
sklearn.tree._tree.Tree._resize_c  
  File "sklearn\tree\_utils.pyx", line 35, in
```

```
sklearn.tree._utils.safe_realloc
MemoryError: could not allocate 4194304 bytes
```

```
-----
-----
2 fits failed with the following error:
Traceback (most recent call last):
  File "c:\Users\prakhar\anaconda\Lib\site-packages\sklearn\
model_selection\_validation.py", line 895, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "c:\Users\prakhar\anaconda\Lib\site-packages\sklearn\base.py",
line 1474, in wrapper
    return fit_method(estimator, *args, **kwargs)
    ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
  File "c:\Users\prakhar\anaconda\Lib\site-packages\sklearn\ensemble\
_forest.py", line 489, in fit
    trees = Parallel(
    ^^^^^^^^^
  File "c:\Users\prakhar\anaconda\Lib\site-packages\sklearn\utils\
parallel.py", line 67, in __call__
    return super().__call__(iterable_with_config)
    ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
  File "c:\Users\prakhar\anaconda\Lib\site-packages\joblib\
parallel.py", line 1918, in __call__
    return output if self.return_generator else list(output)
    ^^^^^^^^^^^^^
  File "c:\Users\prakhar\anaconda\Lib\site-packages\joblib\
parallel.py", line 1847, in _get_sequential_output
    res = func(*args, **kwargs)
    ^^^^^^^^^^^^^^^^^^^^^^^^^
  File "c:\Users\prakhar\anaconda\Lib\site-packages\sklearn\utils\
parallel.py", line 129, in __call__
    return self.function(*args, **kwargs)
    ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
  File "c:\Users\prakhar\anaconda\Lib\site-packages\sklearn\ensemble\
_forest.py", line 192, in _parallel_build_trees
    tree._fit(
  File "c:\Users\prakhar\anaconda\Lib\site-packages\sklearn\tree\
_classes.py", line 472, in _fit
    builder.build(self.tree_, X, y, sample_weight,
missing_values_in_feature_mask)
  File "sklearn\tree\_tree.pyx", line 166, in
sklearn.tree._tree.DepthFirstTreeBuilder.build
  File "sklearn\tree\_tree.pyx", line 285, in
sklearn.tree._tree.DepthFirstTreeBuilder.build
  File "sklearn\tree\_tree.pyx", line 940, in
sklearn.tree._tree.Tree._add_node
  File "sklearn\tree\_tree.pyx", line 908, in
sklearn.tree._tree.Tree._resize_c
```

```
File "sklearn\tree\_utils.pyx", line 35, in
sklearn.tree._utils.safe_realloc
MemoryError: could not allocate 4194304 bytes
```

```
-----
-----
2 fits failed with the following error:
Traceback (most recent call last):
  File "c:\Users\prakhar\anaconda\Lib\site-packages\sklearn\
model_selection\_validation.py", line 895, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "c:\Users\prakhar\anaconda\Lib\site-packages\sklearn\base.py",
line 1474, in wrapper
    return fit_method(estimator, *args, **kwargs)
    ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
  File "c:\Users\prakhar\anaconda\Lib\site-packages\sklearn\ensemble\
_forest.py", line 489, in fit
    trees = Parallel(
    ^^^^^^^^^
  File "c:\Users\prakhar\anaconda\Lib\site-packages\sklearn\utils\
parallel.py", line 67, in __call__
    return super().__call__(iterable_with_config)
    ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
  File "c:\Users\prakhar\anaconda\Lib\site-packages\joblib\
parallel.py", line 1918, in __call__
    return output if self.return_generator else list(output)
    ^^^^^^^^^^^^^
  File "c:\Users\prakhar\anaconda\Lib\site-packages\joblib\
parallel.py", line 1847, in _get_sequential_output
    res = func(*args, **kwargs)
    ^^^^^^^^^^^^^^^^^^^^^^^^^
  File "c:\Users\prakhar\anaconda\Lib\site-packages\sklearn\utils\
parallel.py", line 129, in __call__
    return self.function(*args, **kwargs)
    ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
  File "c:\Users\prakhar\anaconda\Lib\site-packages\sklearn\ensemble\
_forest.py", line 192, in _parallel_build_trees
    tree._fit(
  File "c:\Users\prakhar\anaconda\Lib\site-packages\sklearn\tree\
_classes.py", line 472, in _fit
    builder.build(self.tree_, X, y, sample_weight,
missing_values_in_feature_mask)
  File "sklearn\tree\_tree.pyx", line 166, in
sklearn.tree._tree.DepthFirstTreeBuilder.build
  File "sklearn\tree\_tree.pyx", line 285, in
sklearn.tree._tree.DepthFirstTreeBuilder.build
  File "sklearn\tree\_tree.pyx", line 940, in
sklearn.tree._tree.Tree._add_node
  File "sklearn\tree\_tree.pyx", line 908, in
```

```
sklearn.tree._tree.Tree._resize_c
File "sklearn\tree\_utils.pyx", line 35, in
sklearn.tree._utils.safe_realloc
MemoryError: could not allocate 2097152 bytes
```

```
warnings.warn(some_fits_failed_message, FitFailedWarning)
c:\Users\prakhar\anaconda\Lib\site-packages\sklearn\model_selection\
_search.py:1051: UserWarning: One or more of the test scores are non-
finite: [0.64090456 0.63904023 0.63827256 0.63813402 0.63919608
0.63032461
0.63783967 0.63101147 0.63983675          nan          nan 0.63814557
0.63013413 0.63075173 0.63793778 0.63948467 0.63819752 0.64010226
0.6413663 0.63897674]
warnings.warn(
```

```
-----
-----
KeyboardInterrupt                                Traceback (most recent call
last)
Cell In[45], line 11
      4 for name, model, params in random_cv_model:
      5     random = RandomizedSearchCV(estimator=model,
      6                                 param_distributions=params,
      7                                 n_iter=20,
      8                                 cv=5,
      9                                 verbose=2,
     10                                 n_jobs=-1)
--> 11     random.fit(X_train, Y_train)
     12     model_param[name] = random.best_params_
     14 for model_name in model_param:
```

```
File c:\Users\prakhar\anaconda\Lib\site-packages\sklearn\base.py:1474,
in _fit_context.<locals>.decorator.<locals>.wrapper(estimator, *args,
**kwargs)
     1467     estimator._validate_params()
     1469     with config_context(
     1470         skip_parameter_validation=(
     1471             prefer_skip_nested_validation or
global_skip_validation
     1472         )
     1473     ):
-> 1474     return fit_method(estimator, *args, **kwargs)
```

```
File c:\Users\prakhar\anaconda\Lib\site-packages\sklearn\
model_selection\_search.py:1008, in BaseSearchCV.fit(self, X, y,
**params)
     1006     refit_start_time = time.time()
     1007     if y is not None:
-> 1008         self.best_estimator_.fit(X, y,
**routed_params.estimator.fit)
```

```
1009 else:
1010     self.best_estimator_.fit(X, **routed_params.estimator.fit)
```

File c:\Users\prakhar\anaconda\Lib\site-packages\sklearn\base.py:1474, in \_fit\_context.<locals>.decorator.<locals>.wrapper(estimator, \*args, \*\*kwargs)

```
1467     estimator._validate_params()
1469 with config_context(
1470     skip_parameter_validation=(
1471         prefer_skip_nested_validation or
global_skip_validation
1472     )
1473 ):
-> 1474     return fit_method(estimator, *args, **kwargs)
```

File c:\Users\prakhar\anaconda\Lib\site-packages\sklearn\ensemble\\_forest.py:489, in BaseForest.fit(self, X, y, sample\_weight)

```
478 trees = [
479     self._make_estimator(append=False,
random_state=random_state)
480     for i in range(n_more_estimators)
481 ]
483 # Parallel loop: we prefer the threading backend as the Cython
code
484 # for fitting the trees is internally releasing the Python GIL
485 # making threading more efficient than multiprocessing in
486 # that case. However, for joblib 0.12+ we respect any
487 # parallel_backend contexts set at a higher level,
488 # since correctness does not rely on using threads.
--> 489 trees = Parallel(
490     n_jobs=self.n_jobs,
491     verbose=self.verbose,
492     prefer="threads",
493 )(
494     delayed(_parallel_build_trees)(
495         t,
496         self.bootstrap,
497         X,
498         y,
499         sample_weight,
500         i,
501         len(trees),
502         verbose=self.verbose,
503         class_weight=self.class_weight,
504         n_samples_bootstrap=n_samples_bootstrap,
505         missing_values_in_feature_mask=missing_values_in_feature_mask,
506     )
507     for i, t in enumerate(trees)
```



```
508 )
510 # Collect newly grown trees
511 self.estimateds_.extend(trees)
```

File c:\Users\prakhar\anaconda\Lib\site-packages\sklearn\utils\parallel.py:67, in Parallel.\_\_call\_\_(self, iterable)

```
62 config = get_config()
63 iterable_with_config = (
64     (_with_config(delayed_func, config), args, kwargs)
65     for delayed_func, args, kwargs in iterable
66 )
--> 67 return super().__call__(iterable_with_config)
```

File c:\Users\prakhar\anaconda\Lib\site-packages\joblib\parallel.py:1918, in Parallel.\_\_call\_\_(self, iterable)

```
1916 output = self._get_sequential_output(iterable)
1917 next(output)
-> 1918 return output if self.return_generator else list(output)
1920 # Let's create an ID that uniquely identifies the current
call. If the
1921 # call is interrupted early and that the same instance is
immediately
1922 # re-used, this id will be used to prevent workers that were
1923 # concurrently finalizing a task from the previous call to run
the
1924 # callback.
1925 with self._lock:
```

File c:\Users\prakhar\anaconda\Lib\site-packages\joblib\parallel.py:1847, in Parallel.\_get\_sequential\_output(self, iterable)

```
1845 self.n_dispatched_batches += 1
1846 self.n_dispatched_tasks += 1
-> 1847 res = func(*args, **kwargs)
1848 self.n_completed_tasks += 1
1849 self.print_progress()
```

File c:\Users\prakhar\anaconda\Lib\site-packages\sklearn\utils\parallel.py:129, in \_FuncWrapper.\_\_call\_\_(self, \*args, \*\*kwargs)

```
127 config = {}
128 with config_context(**config):
--> 129 return self.function(*args, **kwargs)
```

File c:\Users\prakhar\anaconda\Lib\site-packages\sklearn\ensemble\\_forest.py:192, in \_parallel\_build\_trees(tree, bootstrap, X, y, sample\_weight, tree\_idx, n\_trees, verbose, class\_weight, n\_samples\_bootstrap, missing\_values\_in\_feature\_mask)

```
189 elif class_weight == "balanced_subsample":
190     curr_sample_weight *=
compute_sample_weight("balanced", y, indices=indices)
--> 192 tree._fit(
```

```

193         X,
194         y,
195         sample_weight=curr_sample_weight,
196         check_input=False,
197
missing_values_in_feature_mask=missing_values_in_feature_mask,
198     )
199 else:
200     tree._fit(
201         X,
202         y,
203         (...)
204     )
missing_values_in_feature_mask=missing_values_in_feature_mask,
205 )
206 )

```

File c:\Users\prakhar\anaconda\Lib\site-packages\sklearn\tree\\_classes.py:472, in BaseDecisionTree.\_fit(self, X, y, sample\_weight, check\_input, missing\_values\_in\_feature\_mask)

```

461 else:
462     builder = BestFirstTreeBuilder(
463         splitter,
464         min_samples_split,
465         (...)
466     )
467     self.min_impurity_decrease,
468     )
--> 472 builder.build(self.tree_, X, y, sample_weight,
missing_values_in_feature_mask)
474 if self.n_outputs_ == 1 and is_classifier(self):
475     self.n_classes_ = self.n_classes_[0]

```

KeyboardInterrupt:

```

models={

    "Random
Forest":RandomForestClassifier(n_estimators=100,min_samples_split=2,min
n_samples_leaf= 1, max_depth= 30, bootstrap= False),

}
for i in range(len(list(models))):
    model = list(models.values())[i]
    model.fit(X_train, Y_train) # Train model

    # Make predictions
    y_train_pred = model.predict(X_train)
    y_test_pred = model.predict(X_test)

    # Training set performance
    model_train_accuracy = accuracy_score(Y_train, y_train_pred) #

```

```

Calculate Accuracy
    model_train_f1 = f1_score(Y_train, y_train_pred,
average='weighted') # Calculate F1-score
    model_train_precision = precision_score(Y_train, y_train_pred) #
Calculate Precision
    model_train_recall = recall_score(Y_train, y_train_pred) #
Calculate Recall
    model_train_rocauc_score = roc_auc_score(Y_train, y_train_pred)

    # Test set performance
    model_test_accuracy = accuracy_score(Y_test, y_test_pred) #
Calculate Accuracy
    model_test_f1 = f1_score(Y_test, y_test_pred, average='weighted')
# Calculate F1-score
    model_test_precision = precision_score(Y_test, y_test_pred) #
Calculate Precision
    model_test_recall = recall_score(Y_test, y_test_pred) # Calculate
Recall
    model_test_rocauc_score = roc_auc_score(Y_test, y_test_pred)
#Calculate Roc

print(list(models.keys())[i])

print('Model performance for Training set')
print("- Accuracy: {:.4f}".format(model_train_accuracy))
print('- F1 score: {:.4f}'.format(model_train_f1))

print('- Precision: {:.4f}'.format(model_train_precision))
print('- Recall: {:.4f}'.format(model_train_recall))
print('- Roc Auc Score: {:.4f}'.format(model_train_rocauc_score))

print('-----')

print('Model performance for Test set')
print('- Accuracy: {:.4f}'.format(model_test_accuracy))
print('- F1 score: {:.4f}'.format(model_test_f1))
print('- Precision: {:.4f}'.format(model_test_precision))
print('- Recall: {:.4f}'.format(model_test_recall))
print('- Roc Auc Score: {:.4f}'.format(model_test_rocauc_score))

print('='*35)
print('\n')

```

Random Forest

Model performance for Training set

- Accuracy: 1.0000
- F1 score: 1.0000
- Precision: 1.0000
- Recall: 1.0000
- Roc Auc Score: 1.0000

-----

Model performance for Test set

- Accuracy: 0.7858
- F1 score: 0.7468
- Precision: 0.8011
- Recall: 0.9565
- Roc Auc Score: 0.5973

=====

```
from sklearn.metrics import confusion_matrix
confusion_matrix(Y_test,)


## Plot ROC AUC Curve
from sklearn.metrics import roc_auc_score, roc_curve
plt.figure()

# Add the models to the list that you want to view on the ROC plot
auc_models = [
    {
        'label': 'Random Forest Classifier',
        'model':
RandomForestClassifier(n_estimators=1000,min_samples_split=2,
max_features=7,max_depth=None),
        'auc': 0.5973
    },
]

# create loop through all model
for algo in auc_models:
    model = algo['model'] # select the model
    model.fit(X_train, Y_train) # train the model
# Compute False positive rate, and True positive rate
    fpr, tpr, thresholds = roc_curve(Y_test,
model.predict_proba(X_test)[:,-1])
# Calculate Area under the curve to display on the plot
    plt.plot(fpr, tpr, label='%s ROC (area = %0.2f)' % (algo['label'],
algo['auc']))

# Custom settings for the plot
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
```

```
plt.xlabel('1-Specificity(False Positive Rate)')
plt.ylabel('Sensitivity(True Positive Rate)')
plt.title('Receiver Operating Characteristic')
plt.legend(loc="lower right")
plt.savefig("auc.png")
plt.show()
```



The achieved ROC AUC score of 0.6 indicates that the model performs moderately better.